

```
julia> using TimeSeries
```

Time and Data Analysis

In data analysis, a time dimension isn't always
necessary

Nobody considers the date and time of day that
Edgar Anderson recorded measurements on
three species of irises

Yet data analysts have still benefited from the 150
timeless observations

Often though, the time dimension is an important piece of information

It provides a distance between observations

TimeSeries is a lightweight package for data
analysis when time matters

Package Design

parsimonious code base

```
[src (master)]
➔ git ls-files | xargs wc -l
    4 .timeseriesrc.jl
   23 TimeSeries.jl
  150 apply.jl
   74 combine.jl
   23 readwrite.jl
   58 split.jl
  203 timearray.jl
  535 total
```

TimeArray is an immutable type to work with
time series data

```
timestamp::Vector{Date}  
values::Array{T,N}  
colnames::Vector{UTF8String}
```

timestamp length matches size of values rows
colnames length matches size of values column
dates cannot be duplicates
dates cannot be in a random unsorted order

The constructor also orders the rows starting with the first observation appearing first

API tour


```
julia> using MarketData
```

```
julia> ohlc
```

```
500x4 TimeArray{Float64,2} 2000-01-03 to 2001-12-31
```

	Open	High	Low	Close
2000-01-03	104.88	112.5	101.69	111.94
2000-01-04	108.25	110.62	101.19	102.5
2000-01-05	103.75	110.56	103.0	104.0
2000-01-06	106.12	107.0	95.0	95.0
:				
2001-12-26	21.35	22.3	21.14	21.49
2001-12-27	21.58	22.25	21.58	22.07
2001-12-28	21.97	23.0	21.96	22.43
2001-12-31	22.51	22.66	21.83	21.9

Columns are indexed by string

```
julia> ohlc["Close"]  
500x1 TimeArray{Float64,1} 2000-01-03 to 2001-12-31
```

	Close
2000-01-03	111.94
2000-01-04	102.5
2000-01-05	104.0
2000-01-06	95.0
⋮	
2001-12-26	21.49
2001-12-27	22.07
2001-12-28	22.43
2001-12-31	21.9

Rows are indexed by date
and date ranges

```
julia> ohlc[Date(2000,1,10)]  
1x4 TimeArray{Float64,2} 2000-01-10 to 2000-01-10
```

	Open	High	Low	Close
2000-01-10	102.0	102.25	94.75	97.75

```
julia> ohlc[Date(2000,1,1):Date(2000,1,10)]  
6x4 TimeArray{Float64,2} 2000-01-03 to 2000-01-10
```

	Open	High	Low	Close
2000-01-03	104.88	112.5	101.69	111.94
2000-01-04	108.25	110.62	101.19	102.5
2000-01-05	103.75	110.56	103.0	104.0
2000-01-06	106.12	107.0	95.0	95.0
2000-01-07	96.5	101.0	95.5	99.5
2000-01-10	102.0	102.25	94.75	97.75

Rows are indexed by integer
and integer ranges

```
julia> ohlc[1]
```

```
1x4 TimeArray{Float64,2} 2000-01-03 to 2000-01-03
```

	Open	High	Low	Close
2000-01-03	104.88	112.5	101.69	111.94

```
julia> ohlc[1:3]
```

```
3x4 TimeArray{Float64,2} 2000-01-03 to 2000-01-05
```

	Open	High	Low	Close
2000-01-03	104.88	112.5	101.69	111.94
2000-01-04	108.25	110.62	101.19	102.5
2000-01-05	103.75	110.56	103.0	104.0

Columns and rows can be
simultaneously indexed


```
julia> ohlc["Open", "Close"][[1:3,12]]  
4x2 TimeArray{Float64,2} 2000-01-03 to 2000-01-19
```

	Open	Close
2000-01-03	104.88	111.94
2000-01-04	108.25	102.5
2000-01-05	103.75	104.0
2000-01-19	105.62	106.56

methods to subset between
specific dates

```
julia> from(ohlc, 2001,12,27)
```

```
3x4 TimeArray{Float64,2} 2001-12-27 to 2001-12-31
```

	Open	High	Low	Close
2001-12-27	21.58	22.25	21.58	22.07
2001-12-28	21.97	23.0	21.96	22.43
2001-12-31	22.51	22.66	21.83	21.9

```
julia> to(ohlc, 2000,1,5)
```

```
3x4 TimeArray{Float64,2} 2000-01-03 to 2000-01-05
```

	Open	High	Low	Close
2000-01-03	104.88	112.5	101.69	111.94
2000-01-04	108.25	110.62	101.19	102.5
2000-01-05	103.75	110.56	103.0	104.0

aggregate on specific time period

```
julia> by(ohlc, 1, period=dayofweek) # Mondays  
95x4 TimeArray{Float64,2} 2000-01-03 to 2001-12-31
```

	Open	High	Low	Close
2000-01-03	104.88	112.5	101.69	111.94
2000-01-10	102.0	102.25	94.75	97.75
2000-01-24	108.44	112.75	105.12	106.25
2000-01-31	101.0	103.88	94.5	103.75
:				
2001-12-10	22.29	22.99	22.23	22.54
2001-12-17	20.4	21.0	20.19	20.62
2001-12-24	20.9	21.45	20.9	21.36
2001-12-31	22.51	22.66	21.83	21.9

collapse allows control on what
values are aggregated to
larger time frame

```
julia> collapse(ohlc["Close"], last, period=month)
24x1 TimeArray{Float64,1} 2000-01-31 to 2001-12-31
```

	Close
2000-01-31	103.75
2000-02-29	114.62
2000-03-31	135.81
2000-04-28	124.06
⋮	
2001-09-28	15.51
2001-10-31	17.56
2001-11-30	21.3
2001-12-31	21.9

element-wise operations


```
julia> ohlc["High"] .- ohlc["Close"]  
500x1 TimeArray{Float64,1} 2000-01-03 to 2001-12-31
```

```
           Hi.-Cl  
2000-01-03 | 0.56  
2000-01-04 | 8.12  
2000-01-05 | 6.56  
2000-01-06 | 12.0  
⋮  
2001-12-26 | 0.81  
2001-12-27 | 0.18  
2001-12-28 | 0.57  
2001-12-31 | 0.76
```

```
julia> ohlc["Open"] .> ohlc["Close"]  
500x1 TimeArray{Bool,1} 2000-01-03 to 2001-12-31
```

```
      Op.>Cl  
2000-01-03 | false  
2000-01-04 | true  
2000-01-05 | false  
2000-01-06 | true  
⋮  
2001-12-26 | false  
2001-12-27 | false  
2001-12-28 | false  
2001-12-31 | true
```

find when a condition is met

```
julia> greendays = findwhen(ohlc["Close"] .> ohlc["Open"]);
```

```
julia> typeof(greendays)  
Array{Date,1}
```

```
julia> ohlc[greendays]  
244x4 TimeArray{Float64,2} 2000-01-03 to 2001-12-28
```

	Open	High	Low	Close
2000-01-03	104.88	112.5	101.69	111.94
2000-01-05	103.75	110.56	103.0	104.0
2000-01-07	96.5	101.0	95.5	99.5
2000-01-13	94.48	98.75	92.5	96.75
:				
2001-12-24	20.9	21.45	20.9	21.36
2001-12-26	21.35	22.3	21.14	21.49
2001-12-27	21.58	22.25	21.58	22.07
2001-12-28	21.97	23.0	21.96	22.43

```
julia> reddays = findall(ohlc["Close"] .< ohlc["Open"]);
```

```
julia> typeof(reddays)
```

```
Array{Int64,1}
```

```
julia> ohlc[reddays]
```

```
252x4 TimeArray{Float64,2} 2000-01-04 to 2001-12-31
```

	Open	High	Low	Close
2000-01-04	108.25	110.62	101.19	102.5
2000-01-06	106.12	107.0	95.0	95.0
2000-01-10	102.0	102.25	94.75	97.75
2000-01-11	95.94	99.38	90.5	92.75
:				
2001-12-14	20.73	20.83	20.09	20.39
2001-12-20	21.4	21.47	20.62	20.67
2001-12-21	21.01	21.54	20.8	21.0
2001-12-31	22.51	22.66	21.83	21.9

time-related transformations

```
julia> lag(ohlc["Open"])  
499x1 TimeArray{Float64,1} 2000-01-04 to 2001-12-31
```

	Open
2000-01-04	104.88
2000-01-05	108.25
2000-01-06	103.75
2000-01-07	106.12
⋮	
2001-12-26	20.9
2001-12-27	21.35
2001-12-28	21.58
2001-12-31	21.97

```
julia> percentchange(ohlc["Close"], method="log")  
499x1 TimeArray{Float64,1} 2000-01-04 to 2001-12-31
```

	Close
2000-01-04	-0.09
2000-01-05	0.01
2000-01-06	-0.09
2000-01-07	0.05
⋮	
2001-12-26	0.01
2001-12-27	0.03
2001-12-28	0.02
2001-12-31	-0.02


```
julia> moving(ohlc["Close"], mean, 20)
481x1 TimeArray{Float64,1} 2000-01-31 to 2001-12-31
```

	Close
2000-01-31	103.36
2000-02-01	102.78
2000-02-02	102.59
2000-02-03	102.56
⋮	
2001-12-26	21.49
2001-12-27	21.56
2001-12-28	21.67
2001-12-31	21.7

```
julia> upto(ohl["Close"], sum)
500x1 TimeArray{Float64,1} 2000-01-03 to 2001-12-31
```

	Close
2000-01-03	111.94
2000-01-04	214.44
2000-01-05	318.44
2000-01-06	413.44
:	
2001-12-26	23028.84
2001-12-27	23050.91
2001-12-28	23073.34
2001-12-31	23095.24

basecall uses fast base methods

```
julia> BA["Close"]  
13090x1 TimeArray{Float64,1} 1962-01-02 to 2013-12-31
```

	Close
1962-01-02	50.0
1962-01-03	51.0
1962-01-04	50.5
1962-01-05	49.5
:	
2013-12-26	138.27
2013-12-27	136.9
2013-12-30	135.92
2013-12-31	136.49

```
julia> @time upto(BA["Close"], sum);  
elapsed time: 0.092170981 seconds (4663992 bytes allocated)
```

```
julia> @time basecall(BA["Close"], cumsum);  
elapsed time: 0.0099391 seconds (3990200 bytes allocated)
```

merging two TimeArrays

```
julia> @time merge(BA["High"], CAT["Low"])
elapsed time: 1.776906707 seconds (12571880 bytes allocated)
13090x2 TimeArray{Float64,2} 1962-01-02 to 2013-12-31
```


	High	Low
1962-01-02	50.88	38.12
1962-01-03	51.75	38.12
1962-01-04	51.88	39.75
1962-01-05	50.75	39.75
:		
2013-12-26	138.59	90.7
2013-12-27	138.88	90.56
2013-12-30	137.37	90.28
2013-12-31	137.05	90.46

merge is slow and colnames are not unique.
Pull requests are welcome!

When Observations
Disappear

missing values fall into black holes

This is what happens when your data disappears



```
julia> ohlc[1]
1x4 TimeArray{Float64,2} 2000-01-03 to 2000-01-03
```

	Open	High	Low	Close
2000-01-03	104.88	112.5	101.69	111.94

```
julia> lag(ohlc)
499x4 TimeArray{Float64,2} 2000-01-04 to 2001-12-31
```

	Open	High	Low	Close
2000-01-04	104.88	112.5	101.69	111.94
2000-01-05	108.25	110.62	101.19	102.5
2000-01-06	103.75	110.56	103.0	104.0
2000-01-07	106.12	107.0	95.0	95.0
:				
2001-12-26	20.9	21.45	20.9	21.36
2001-12-27	21.35	22.3	21.14	21.49
2001-12-28	21.58	22.25	21.58	22.07
2001-12-31	21.97	23.0	21.96	22.43

Rather than conflate NaN with missingness
or implement the DataFrames NA type,
TimeSeries tosses consumed values
into black holes

This is not ideal, and other packages work around this by introducing NaN placeholders similar to how pandas handles this

```
julia> using Quandl
```

```
julia> quandl("CHRIS/CME_DK1") # Class IV Milk Futures  
100x8 TimeArray{Float64,2} 2014-01-14 to 2014-06-05
```

	Open	High	Low	Last	Change	Settle	Volume	Open Interest
2014-01-14	NaN	22.05	22.02	NaN	NaN	22.05	NaN	NaN
2014-01-15	NaN	22.1	21.99	NaN	NaN	22.09	NaN	NaN
2014-01-16	NaN	22.2	22.1	NaN	NaN	22.1	NaN	NaN
2014-01-17	NaN	22.2	22.17	NaN	NaN	22.2	NaN	NaN
:								
2014-06-02	NaN	NaN	NaN	NaN	NaN	22.59	0.0	1660.0
2014-06-03	NaN	NaN	22.54	NaN	0.03	22.56	0.0	1660.0
2014-06-04	22.55	22.65	22.55	NaN	NaN	22.55	11.0	1565.0
2014-06-05	22.57	22.7	22.57	NaN	0.14	22.69	6.0	1567.0

Time Series in R and Python

R's xts and Python's pandas
for comparison

R code to duplicate data object

```
> library(quantmod)
> getSymbols('AAPL', from='2000-1-1', to='2001-12-31')
[1] "AAPL"
> head(AAPL)
```

	AAPL.Open	AAPL.High	AAPL.Low	AAPL.Close	AAPL.Volume	AAPL.Adjusted
2000-01-03	104.88	112.50	101.69	111.94	133949200	3.82
2000-01-04	108.25	110.62	101.19	102.50	128094400	3.50
2000-01-05	103.75	110.56	103.00	104.00	194580400	3.55
2000-01-06	106.13	107.00	95.00	95.00	191993200	3.24
2000-01-07	96.50	101.00	95.50	99.50	115183600	3.40
2000-01-10	102.00	102.25	94.75	97.75	126266000	3.34

pandas code to duplicate data object

```
In [1]: from pandas import *
```

```
In [2]: from pandas.io.data import DataReader
```

```
In [3]: from datetime import datetime
```

```
In [4]: AAPL = DataReader("AAPL", "yahoo", datetime(2000,1,1), datetime(2001,12,31))
```

```
In [5]: AAPL[0:3]
```

```
Out[5]:
```

	Open	High	Low	Close	Volume	Adj Close
Date						
2000-01-03	104.88	112.50	101.69	111.94	133949200	3.82
2000-01-04	108.25	110.62	101.19	102.50	128094400	3.50
2000-01-05	103.75	110.56	103.00	104.00	194580400	3.55

Both xts and pandas support date strings for
subsetting data

This is a convenience that TimeSeries
is taking a pass on

R's xts doesn't do error-checking that your
string is ISO-compliant and fails
silently when it isn't

Both xts and pandas preserve dates (rows) of data consumed by shifting or calculation


```
> lagged = lag(AAPL)
```

```
> lagged[1:3]
```

	AAPL.Open	AAPL.High	AAPL.Low	AAPL.Close	AAPL.Volume	AAPL.Adjusted
2000-01-03	NA	NA	NA	NA	NA	NA
2000-01-04	104.88	112.50	101.69	111.94	133949200	3.82
2000-01-05	108.25	110.62	101.19	102.50	128094400	3.50

```
In [5]: lag = AAPL.shift(1)
```

```
In [6]: lag[0:3]
```

```
Out[6]:
```

	Open	High	Low	Close	Volume	Adj Close
Date						
2000-01-03	NaN	NaN	NaN	NaN	NaN	NaN
2000-01-04	104.88	112.50	101.69	111.94	133949200	3.82
2000-01-05	108.25	110.62	101.19	102.50	128094400	3.50

Both xts and pandas show floats disguised as integers, when they were converted to floats from integers to fit into their array

TimeSeries defaults to showing
all elements of a float array
as floats, but you can modify that

You can also modify how many
decimals are displayed

.timeseriesrc.jl

```
##### customizable show #####
```

```
const DECIMALS = 4    # default value is 2
```

```
const SHOWINT = true  # defaults to false
```

```
julia> ohlcv[1:3]
```

```
3x5 TimeArray{Float64,2} 2000-01-03 to 2000-01-05
```

	Open	High	Low	Close	Volume
2000-01-03	104.88	112.5	101.69	111.94	4783900
2000-01-04	108.25	110.62	101.19	102.5	4574800
2000-01-05	103.75	110.56	103.0	104.0	6949300

```
julia> percentchange(cl)[1:3]
```

```
3x1 TimeArray{Float64,2} 2000-01-04 to 2000-01-06
```

	Close
2000-01-04	-0.0843
2000-01-05	0.0146
2000-01-06	-0.0865

TimeSeries lives in the JuliaStats organization

Special thanks to Jacob Quinn for his work on the Dates.jl package. TimeSeries now depends on that package.

The End.

Time for a break.