

```
julia> using TimeSeries
```

Time and Data Analysis

In data analysis, a time dimension isn't always
necessary

Nobody recorded when Sir Ronald Fisher measured iris petal and sepal lengths and widths

And yet researchers, professors and their students have still benefited from the 150 timeless observations

Often though, the time dimension is an important
piece of information

It provides a distance (though imperfect) between
observations

TimeSeries is a lightweight package for data analysis when time matters

Package Design

parsimonious code base


```
[src (master)]
```

```
✈ git ls-files | xargs wc -l  
    4 .timeseriesrc.jl  
   23 TimeSeries.jl  
  150 apply.jl  
   74 combine.jl  
   23 readwrite.jl  
   58 split.jl  
  203 timearray.jl  
  535 total
```

TimeArray is an immutable type to work with
time series data

It includes the following elements:

timestamp::Vector{Date}

values::Array{T,N}

colnames::Vector{UTF8String}

Inner constructor enforces 4 invariants:

timestamp and values must have same row length

dates cannot be duplicated

dates cannot be in a random unsorted order

colnames must match the width of values array

The constructor also orients the rows starting with oldest and moving towards newest observations

API tour

```
julia> using MarketData
```

```
julia> ohlc
```

```
500x4 TimeArray{Float64,2} 2000-01-03 to 2001-12-31
```

| | | Open | High | Low | Close |
|------------|--|--------|--------|--------|--------|
| 2000-01-03 | | 104.88 | 112.5 | 101.69 | 111.94 |
| 2000-01-04 | | 108.25 | 110.62 | 101.19 | 102.5 |
| 2000-01-05 | | 103.75 | 110.56 | 103.0 | 104.0 |
| 2000-01-06 | | 106.12 | 107.0 | 95.0 | 95.0 |
| ⋮ | | | | | |
| 2001-12-26 | | 21.35 | 22.3 | 21.14 | 21.49 |
| 2001-12-27 | | 21.58 | 22.25 | 21.58 | 22.07 |
| 2001-12-28 | | 21.97 | 23.0 | 21.96 | 22.43 |
| 2001-12-31 | | 22.51 | 22.66 | 21.83 | 21.9 |

Columns are indexed by string


```
julia> ohlc["Close"]
```

```
500x1 TimeArray{Float64,1} 2000-01-03 to 2001-12-31
```

| | Close |
|------------|--------|
| 2000-01-03 | 111.94 |
| 2000-01-04 | 102.5 |
| 2000-01-05 | 104.0 |
| 2000-01-06 | 95.0 |
| : | |
| 2001-12-26 | 21.49 |
| 2001-12-27 | 22.07 |
| 2001-12-28 | 22.43 |
| 2001-12-31 | 21.9 |

Rows are indexed by date
and date ranges

```
julia> ohlc[Date(2000,1,10)]
```

```
1x4 TimeArray{Float64,2} 2000-01-10 to 2000-01-10
```

| | Open | High | Low | Close |
|------------|-------|--------|-------|-------|
| 2000-01-10 | 102.0 | 102.25 | 94.75 | 97.75 |

```
julia> ohlc[Date(2000,1,1):Date(2000,1,10)]
```

```
6x4 TimeArray{Float64,2} 2000-01-03 to 2000-01-10
```

| | Open | High | Low | Close |
|------------|--------|--------|--------|--------|
| 2000-01-03 | 104.88 | 112.5 | 101.69 | 111.94 |
| 2000-01-04 | 108.25 | 110.62 | 101.19 | 102.5 |
| 2000-01-05 | 103.75 | 110.56 | 103.0 | 104.0 |
| 2000-01-06 | 106.12 | 107.0 | 95.0 | 95.0 |
| 2000-01-07 | 96.5 | 101.0 | 95.5 | 99.5 |
| 2000-01-10 | 102.0 | 102.25 | 94.75 | 97.75 |

Rows are indexed by integer
and integer ranges

```
julia> ohlc[1]
```

```
1x4 TimeArray{Float64,2} 2000-01-03 to 2000-01-03
```

| | Open | High | Low | Close |
|------------|--------|-------|--------|--------|
| 2000-01-03 | 104.88 | 112.5 | 101.69 | 111.94 |

```
julia> ohlc[1:3]
```

```
3x4 TimeArray{Float64,2} 2000-01-03 to 2000-01-05
```

| | Open | High | Low | Close |
|------------|--------|--------|--------|--------|
| 2000-01-03 | 104.88 | 112.5 | 101.69 | 111.94 |
| 2000-01-04 | 108.25 | 110.62 | 101.19 | 102.5 |
| 2000-01-05 | 103.75 | 110.56 | 103.0 | 104.0 |

Columns and rows can be
simultaneously indexed

```
julia> ohlc["Open", "Close"][[1:3,12]]  
4x2 TimeArray{Float64,2} 2000-01-03 to 2000-01-19
```

| | Open | Close |
|------------|--------|--------|
| 2000-01-03 | 104.88 | 111.94 |
| 2000-01-04 | 108.25 | 102.5 |
| 2000-01-05 | 103.75 | 104.0 |
| 2000-01-19 | 105.62 | 106.56 |

methods to subset between
specific dates


```
julia> from(ohlc, 2001,12,27)
```

```
3x4 TimeArray{Float64,2} 2001-12-27 to 2001-12-31
```

| | | Open | High | Low | Close |
|------------|--|-------|-------|-------|-------|
| 2001-12-27 | | 21.58 | 22.25 | 21.58 | 22.07 |
| 2001-12-28 | | 21.97 | 23.0 | 21.96 | 22.43 |
| 2001-12-31 | | 22.51 | 22.66 | 21.83 | 21.9 |

```
julia> to(ohlc, 2000,1,5)
```

```
3x4 TimeArray{Float64,2} 2000-01-03 to 2000-01-05
```

| | | Open | High | Low | Close |
|------------|--|--------|--------|--------|--------|
| 2000-01-03 | | 104.88 | 112.5 | 101.69 | 111.94 |
| 2000-01-04 | | 108.25 | 110.62 | 101.19 | 102.5 |
| 2000-01-05 | | 103.75 | 110.56 | 103.0 | 104.0 |

aggregate on specific time period

```
julia> by(ohlc, 1, period=dayofweek) # Mondays
95x4 TimeArray{Float64,2} 2000-01-03 to 2001-12-31
```

| | Open | High | Low | Close |
|------------|--------|--------|--------|--------|
| 2000-01-03 | 104.88 | 112.5 | 101.69 | 111.94 |
| 2000-01-10 | 102.0 | 102.25 | 94.75 | 97.75 |
| 2000-01-24 | 108.44 | 112.75 | 105.12 | 106.25 |
| 2000-01-31 | 101.0 | 103.88 | 94.5 | 103.75 |
| ⋮ | | | | |
| 2001-12-10 | 22.29 | 22.99 | 22.23 | 22.54 |
| 2001-12-17 | 20.4 | 21.0 | 20.19 | 20.62 |
| 2001-12-24 | 20.9 | 21.45 | 20.9 | 21.36 |
| 2001-12-31 | 22.51 | 22.66 | 21.83 | 21.9 |

collapse allows control on what
values are aggregated to
larger time frame

```
julia> collapse(ohlc["Close"], last, period=month)
24x1 TimeArray{Float64,1} 2000-01-31 to 2001-12-31
```

| | Close |
|------------|--------|
| 2000-01-31 | 103.75 |
| 2000-02-29 | 114.62 |
| 2000-03-31 | 135.81 |
| 2000-04-28 | 124.06 |
| : | |
| 2001-09-28 | 15.51 |
| 2001-10-31 | 17.56 |
| 2001-11-30 | 21.3 |
| 2001-12-31 | 21.9 |

element-wise operations

```
julia> ohlc["High"] .- ohlc["Close"]  
500x1 TimeArray{Float64,1} 2000-01-03 to 2001-12-31
```

| | Hi.-Cl |
|------------|--------|
| 2000-01-03 | 0.56 |
| 2000-01-04 | 8.12 |
| 2000-01-05 | 6.56 |
| 2000-01-06 | 12.0 |
| ⋮ | |
| 2001-12-26 | 0.81 |
| 2001-12-27 | 0.18 |
| 2001-12-28 | 0.57 |
| 2001-12-31 | 0.76 |

```
julia> ohlc["Open"] .> ohlc["Close"]  
500x1 TimeArray{Bool,1} 2000-01-03 to 2001-12-31
```

| | Op.>Cl |
|------------|--------|
| 2000-01-03 | false |
| 2000-01-04 | true |
| 2000-01-05 | false |
| 2000-01-06 | true |
| ⋮ | |
| 2001-12-26 | false |
| 2001-12-27 | false |
| 2001-12-28 | false |
| 2001-12-31 | true |

find when a condition is met

```
julia> greendays = findwhen(ohlc["Close"] .> ohlc["Open"]);
```

```
julia> typeof(greendays)
```

```
Array{Date,1}
```

```
julia> ohlc[greendays]
```

```
244x4 TimeArray{Float64,2} 2000-01-03 to 2001-12-28
```

| | Open | High | Low | Close |
|------------|--------|--------|--------|--------|
| 2000-01-03 | 104.88 | 112.5 | 101.69 | 111.94 |
| 2000-01-05 | 103.75 | 110.56 | 103.0 | 104.0 |
| 2000-01-07 | 96.5 | 101.0 | 95.5 | 99.5 |
| 2000-01-13 | 94.48 | 98.75 | 92.5 | 96.75 |
| ⋮ | | | | |
| 2001-12-24 | 20.9 | 21.45 | 20.9 | 21.36 |
| 2001-12-26 | 21.35 | 22.3 | 21.14 | 21.49 |
| 2001-12-27 | 21.58 | 22.25 | 21.58 | 22.07 |
| 2001-12-28 | 21.97 | 23.0 | 21.96 | 22.43 |

```
julia> reddays = findall(ohlc["Close"] .< ohlc["Open"]);
```

```
julia> typeof(reddays)
```

```
Array{Int64,1}
```

```
julia> ohlc[reddays]
```

```
252x4 TimeArray{Float64,2} 2000-01-04 to 2001-12-31
```

| | | Open | High | Low | Close |
|------------|--|--------|--------|--------|-------|
| 2000-01-04 | | 108.25 | 110.62 | 101.19 | 102.5 |
| 2000-01-06 | | 106.12 | 107.0 | 95.0 | 95.0 |
| 2000-01-10 | | 102.0 | 102.25 | 94.75 | 97.75 |
| 2000-01-11 | | 95.94 | 99.38 | 90.5 | 92.75 |
| : | | | | | |
| 2001-12-14 | | 20.73 | 20.83 | 20.09 | 20.39 |
| 2001-12-20 | | 21.4 | 21.47 | 20.62 | 20.67 |
| 2001-12-21 | | 21.01 | 21.54 | 20.8 | 21.0 |
| 2001-12-31 | | 22.51 | 22.66 | 21.83 | 21.9 |

time-related transformations

```
julia> lag(ohlc["Open"])
```

```
499x1 TimeArray{Float64,1} 2000-01-04 to 2001-12-31
```

| | Open |
|------------|--------|
| 2000-01-04 | 104.88 |
| 2000-01-05 | 108.25 |
| 2000-01-06 | 103.75 |
| 2000-01-07 | 106.12 |
| ⋮ | |
| 2001-12-26 | 20.9 |
| 2001-12-27 | 21.35 |
| 2001-12-28 | 21.58 |
| 2001-12-31 | 21.97 |

```
julia> percentchange(ohlc["Close"], method="log")  
499x1 TimeArray{Float64,1} 2000-01-04 to 2001-12-31
```

| | Close |
|------------|-------|
| 2000-01-04 | -0.09 |
| 2000-01-05 | 0.01 |
| 2000-01-06 | -0.09 |
| 2000-01-07 | 0.05 |
| ⋮ | |
| 2001-12-26 | 0.01 |
| 2001-12-27 | 0.03 |
| 2001-12-28 | 0.02 |
| 2001-12-31 | -0.02 |

```
julia> moving(ohlc["Close"], mean, 20)
481x1 TimeArray{Float64,1} 2000-01-31 to 2001-12-31
```

| | Close |
|------------|--------|
| 2000-01-31 | 103.36 |
| 2000-02-01 | 102.78 |
| 2000-02-02 | 102.59 |
| 2000-02-03 | 102.56 |
| ⋮ | |
| 2001-12-26 | 21.49 |
| 2001-12-27 | 21.56 |
| 2001-12-28 | 21.67 |
| 2001-12-31 | 21.7 |


```
julia> upto(ohlc["Close"], sum)
```

```
500x1 TimeArray{Float64,1} 2000-01-03 to 2001-12-31
```

| | Close |
|------------|----------|
| 2000-01-03 | 111.94 |
| 2000-01-04 | 214.44 |
| 2000-01-05 | 318.44 |
| 2000-01-06 | 413.44 |
| ⋮ | |
| 2001-12-26 | 23028.84 |
| 2001-12-27 | 23050.91 |
| 2001-12-28 | 23073.34 |
| 2001-12-31 | 23095.24 |

basecall uses fast base methods

```
julia> BA["Close"]  
13090x1 TimeArray{Float64,1} 1962-01-02 to 2013-12-31
```

| | Close |
|------------|--------|
| 1962-01-02 | 50.0 |
| 1962-01-03 | 51.0 |
| 1962-01-04 | 50.5 |
| 1962-01-05 | 49.5 |
| ⋮ | |
| 2013-12-26 | 138.27 |
| 2013-12-27 | 136.9 |
| 2013-12-30 | 135.92 |
| 2013-12-31 | 136.49 |

```
julia> @time upto(BA["Close"], sum);  
elapsed time: 0.092170981 seconds (4663992 bytes allocated)
```

```
julia> @time basecall(BA["Close"], cumsum);  
elapsed time: 0.0099391 seconds (3990200 bytes allocated)
```

merging two TimeArrays

```
julia> @time merge(BA["High"], CAT["Low"])  
elapsed time: 1.776906707 seconds (12571880 bytes allocated)  
13090x2 TimeArray{Float64,2} 1962-01-02 to 2013-12-31
```


| | High | Low |
|------------|--------|-------|
| 1962-01-02 | 50.88 | 38.12 |
| 1962-01-03 | 51.75 | 38.12 |
| 1962-01-04 | 51.88 | 39.75 |
| 1962-01-05 | 50.75 | 39.75 |
| ⋮ | | |
| 2013-12-26 | 138.59 | 90.7 |
| 2013-12-27 | 138.88 | 90.56 |
| 2013-12-30 | 137.37 | 90.28 |
| 2013-12-31 | 137.05 | 90.46 |

merge is slow and colnames are not unique.
Pull requests are welcome!

When Observations
Disappear

missing values fall into black holes

This is what happens when your data disappears



```
julia> ohlc[1]
```

```
1x4 TimeArray{Float64,2} 2000-01-03 to 2000-01-03
```

| | Open | High | Low | Close |
|------------|--------|-------|--------|--------|
| 2000-01-03 | 104.88 | 112.5 | 101.69 | 111.94 |

```
julia> lag(ohlc)
```

```
499x4 TimeArray{Float64,2} 2000-01-04 to 2001-12-31
```

| | Open | High | Low | Close |
|------------|--------|--------|--------|--------|
| 2000-01-04 | 104.88 | 112.5 | 101.69 | 111.94 |
| 2000-01-05 | 108.25 | 110.62 | 101.19 | 102.5 |
| 2000-01-06 | 103.75 | 110.56 | 103.0 | 104.0 |
| 2000-01-07 | 106.12 | 107.0 | 95.0 | 95.0 |
| ⋮ | | | | |
| 2001-12-26 | 20.9 | 21.45 | 20.9 | 21.36 |
| 2001-12-27 | 21.35 | 22.3 | 21.14 | 21.49 |
| 2001-12-28 | 21.58 | 22.25 | 21.58 | 22.07 |
| 2001-12-31 | 21.97 | 23.0 | 21.96 | 22.43 |

Rather than conflate NaN with missingness
or implement the DataFrames NA type,
TimeSeries tosses consumed values
into black holes

This is not ideal, and other packages work around this by introducing NaN placeholders similar to how pandas handles this

```
julia> using Quandl
```

```
julia> quandl("CHRIS/CME_DK1") # Class IV Milk Futures
```

```
100x8 TimeArray{Float64,2} 2014-01-14 to 2014-06-05
```

| | | Open | High | Low | Last | Change | Settle | Volume | Open Interest |
|------------|--|-------|-------|-------|------|--------|--------|--------|---------------|
| 2014-01-14 | | NaN | 22.05 | 22.02 | NaN | NaN | 22.05 | NaN | NaN |
| 2014-01-15 | | NaN | 22.1 | 21.99 | NaN | NaN | 22.09 | NaN | NaN |
| 2014-01-16 | | NaN | 22.2 | 22.1 | NaN | NaN | 22.1 | NaN | NaN |
| 2014-01-17 | | NaN | 22.2 | 22.17 | NaN | NaN | 22.2 | NaN | NaN |
| ⋮ | | | | | | | | | |
| 2014-06-02 | | NaN | NaN | NaN | NaN | NaN | 22.59 | 0.0 | 1660.0 |
| 2014-06-03 | | NaN | NaN | 22.54 | NaN | 0.03 | 22.56 | 0.0 | 1660.0 |
| 2014-06-04 | | 22.55 | 22.65 | 22.55 | NaN | NaN | 22.55 | 11.0 | 1565.0 |
| 2014-06-05 | | 22.57 | 22.7 | 22.57 | NaN | 0.14 | 22.69 | 6.0 | 1567.0 |

Time Series in R and Python

R's xts and Python's pandas
for comparison

R code to duplicate data object


```
> library(quantmod)
> getSymbols('AAPL', from='2000-1-1', to='2001-12-31')
[1] "AAPL"
> head(AAPL)
```

| | AAPL.Open | AAPL.High | AAPL.Low | AAPL.Close | AAPL.Volume | AAPL.Adjusted |
|------------|-----------|-----------|----------|------------|-------------|---------------|
| 2000-01-03 | 104.88 | 112.50 | 101.69 | 111.94 | 133949200 | 3.82 |
| 2000-01-04 | 108.25 | 110.62 | 101.19 | 102.50 | 128094400 | 3.50 |
| 2000-01-05 | 103.75 | 110.56 | 103.00 | 104.00 | 194580400 | 3.55 |
| 2000-01-06 | 106.13 | 107.00 | 95.00 | 95.00 | 191993200 | 3.24 |
| 2000-01-07 | 96.50 | 101.00 | 95.50 | 99.50 | 115183600 | 3.40 |
| 2000-01-10 | 102.00 | 102.25 | 94.75 | 97.75 | 126266000 | 3.34 |

pandas code to duplicate data object

```
In [1]: from pandas import *
```

```
In [2]: from pandas.io.data import DataReader
```

```
In [3]: from datetime import datetime
```

```
In [4]: AAPL = DataReader("AAPL", "yahoo", datetime(2000,1,1), datetime(2001,12,31))
```

```
In [5]: AAPL[0:3]
```

```
Out[5]:
```

| | Open | High | Low | Close | Volume | Adj Close |
|------------|--------|--------|--------|--------|-----------|-----------|
| Date | | | | | | |
| 2000-01-03 | 104.88 | 112.50 | 101.69 | 111.94 | 133949200 | 3.82 |
| 2000-01-04 | 108.25 | 110.62 | 101.19 | 102.50 | 128094400 | 3.50 |
| 2000-01-05 | 103.75 | 110.56 | 103.00 | 104.00 | 194580400 | 3.55 |

Both xts and pandas support date strings for
subsetting data

This is a convenience that TimeSeries
is taking a pass on

R's xts doesn't do error-checking that your
string is ISO-compliant and fails
silently when it isn't

Both xts and pandas preserve dates (rows) of data consumed by shifting or calculation

```
> lagged = lag(AAPL)
```

```
> lagged[1:3]
```

| | AAPL.Open | AAPL.High | AAPL.Low | AAPL.Close | AAPL.Volume | AAPL.Adjusted |
|------------|-----------|-----------|----------|------------|-------------|---------------|
| 2000-01-03 | NA | NA | NA | NA | NA | NA |
| 2000-01-04 | 104.88 | 112.50 | 101.69 | 111.94 | 133949200 | 3.82 |
| 2000-01-05 | 108.25 | 110.62 | 101.19 | 102.50 | 128094400 | 3.50 |

```
In [5]: lag = AAPL.shift(1)
```

```
In [6]: lag[0:3]
```

```
Out[6]:
```

| | Open | High | Low | Close | Volume | Adj Close |
|------------|--------|--------|--------|--------|-----------|-----------|
| Date | | | | | | |
| 2000-01-03 | NaN | NaN | NaN | NaN | NaN | NaN |
| 2000-01-04 | 104.88 | 112.50 | 101.69 | 111.94 | 133949200 | 3.82 |
| 2000-01-05 | 108.25 | 110.62 | 101.19 | 102.50 | 128094400 | 3.50 |

Both xts and pandas show floats disguised as integers, when they were converted to floats from integers to fit into their array

TimeSeries defaults to showing
all elements of a float array
as floats, but you can modify that

You can also modify how many
decimals are displayed

.timeseriesrc.jl

```
##### customizable show #####
```

```
const DECIMALS = 4    # default value is 2  
const SHOWINT = true  # defaults to false
```

```
julia> ohlcv[1:3]
```

```
3x5 TimeArray{Float64,2} 2000-01-03 to 2000-01-05
```

| | Open | High | Low | Close | Volume |
|------------|--------|--------|--------|--------|---------|
| 2000-01-03 | 104.88 | 112.5 | 101.69 | 111.94 | 4783900 |
| 2000-01-04 | 108.25 | 110.62 | 101.19 | 102.5 | 4574800 |
| 2000-01-05 | 103.75 | 110.56 | 103.0 | 104.0 | 6949300 |

```
julia> percentchange(c1)[1:3]
```

```
3x1 TimeArray{Float64,2} 2000-01-04 to 2000-01-06
```

| | Close |
|------------|---------|
| 2000-01-04 | -0.0843 |
| 2000-01-05 | 0.0146 |
| 2000-01-06 | -0.0865 |

TimeSeries lives in the JuliaStats organization

Special thanks to Jacob Quinn for his work on the Dates.jl package. TimeSeries now depends on that package.

The End.

Time for a break.